

PTRACE How-to

2015/10/8

Wooyeon Lee

ptrace

- Linux built-in process tracing mechanism.
- Provided as a system call.
- With ptrace,
 - a parent process (tracer) can trace its child processes (tracee).
 - can read/write to register and memory of tracee process.

ptrace

- `#include <sys/ptrace.h>`
- `long ptrace(enum __ptrace_request request,
 pid_t pid, void *addr, void *data);`
- => `ptrace(REQ_CODE, PID, ADDR, DATA)`
 - REQ_CODE: a code representing one of following available actions
 - start/end tracing, read/write data, etc.
 - PID: a process id of tracee
 - ADDR: a memory address where you read/write data from/to
 - DATA: a data that you read/write
- Enter 'man ptrace' in your command-line

Start tracing

- `PTRACE_TRACEME`
 - `ptrace(PTRACE_TRACEME, 0, NULL, NULL)`
 - Only `ptrace` method invoked by child
- `PTRACE_ATTACH` / `PTRACE_DETACH`
 - `ptrace(PTRACE_ATTACH, pid, NULL, NULL)`
 - `ptrace(PTRACE_DETACH, pid, NULL, NULL)`
 - It makes a target process as its child

Example

```
pid_t child;  
child = fork();  
  
if (child == 0) { /* child process */  
    ptrace(PTRACE_TRACEME, 0, NULL, NULL);  
    execvp("/bin/ls");  
    // INACCESSIBLE  
}  
else { /* parent process */  
    waitpid(child, ...)  
    // WE GOT CHILD!  
}
```

Register Access

- PTRACE_GETREGS
 - ptrace(PTRACE_GETREGS, pid, NULL, ®s)
- PTRACE_SETREGS
 - ptrace(PTRACE_SETREGS, pid, NULL, ®s)

// data structure representing registers

struct user_regs_struct regs;

```
struct user_regs_struct
{
    __extension__ unsigned long long int r15;
    __extension__ unsigned long long int r14;
    __extension__ unsigned long long int r13;
    __extension__ unsigned long long int r12;
    __extension__ unsigned long long int rbp;
    __extension__ unsigned long long int rbx;
    __extension__ unsigned long long int r11;
    __extension__ unsigned long long int r10;
    __extension__ unsigned long long int r9;
    __extension__ unsigned long long int r8;
    __extension__ unsigned long long int rax;
    __extension__ unsigned long long int rcx;
    __extension__ unsigned long long int rdx;
    __extension__ unsigned long long int rsi;
    __extension__ unsigned long long int rdi;
    __extension__ unsigned long long int orig_rax;
    __extension__ unsigned long long int rip;
    __extension__ unsigned long long int cs;
    __extension__ unsigned long long int eflags;
    __extension__ unsigned long long int rsp;
    __extension__ unsigned long long int ss;
    __extension__ unsigned long long int fs_base;
    __extension__ unsigned long long int gs_base;
    __extension__ unsigned long long int ds;
    __extension__ unsigned long long int es;
    __extension__ unsigned long long int fs;
    __extension__ unsigned long long int gs;
};
```

HAL >> /usr/include/x86_64-linux-gnu/sys/user.h R0

Example

```
#include <sys/user.h>
```

```
struct user_regs_struct regs;
```

```
ptrace(PTRACE_GETREGS, child, NULL, &regs);
```

```
//registers
```

```
unsigned long long int rax = regs.rax;
```

```
unsigned long long int rdi = regs.rdi;
```

```
...
```

System call convention

- System call number & return value

arch/ABI	instruction	syscall #	retval
x86_64	syscall	rax	rax

- System call arguments.

arch/ABI	arg1	arg2	arg3	arg4	arg5	arg6
X86_64	rdi	rsi	rdx	r10	r8	r9

- Enter “man syscall” in command-line

Memory Access

- PTRACE_PEEKDATA/TEXT
 - `data = ptrace(PTRACE_PEEKDATA, pid, addr, NULL)`
- PTRACE_POKEADATA/TEXT
 - `ptrace(PTRACE_POKEADATA, pid, addr, &buf)`
- Transfer(copy) word data.
 - A word contains 64bits for a 64bit architecture.
 - 32bits for 32bit arch

Header Files

- /usr/include/x86_64-linux-gnu/
 - sys/reg.h
 - sys/user.h
 - sys/syscall.h -> bits/syscall.h
 - sys/ptrace.h
 - asm/ptrace-abi.h

PRACTICE

- Tracing an incremental counter
- Manipulating a behavior of program

Optional: USER area Access

- PTRACE_PEEKUSER
- PTRACE_POKEUSER
- USER area holds the registers and other information about the process (<sys/user.h>)

Thanks

`sec-tas@cmslab.snu.ac.kr`

References:

<http://man7.org/linux/man-pages/man2/ptrace.2.html>

[http://www.ee.ryerson.ca/~courses/coe518/LinuxJournal/elj2002-103-
ptrace1.pdf](http://www.ee.ryerson.ca/~courses/coe518/LinuxJournal/elj2002-103-ptrace1.pdf)

[http://delivery.acm.org/10.1145/610000/603626/6210.html?ip=147.4
6.246.161&id=603626&acc=ACTIVE%20SERVICE&key=0EC22F8658578
FE1%2ED83A6478590749B7%2E4D4702B0C3E38B35%2E4D4702B0C3E
38B35&CFID=550559368&CFTOKEN=13078748&__acm__=144420590
9_e1b0402c13a974f603db2271da839e99](http://delivery.acm.org/10.1145/610000/603626/6210.html?ip=147.46.246.161&id=603626&acc=ACTIVE%20SERVICE&key=0EC22F8658578FE1%2ED83A6478590749B7%2E4D4702B0C3E38B35%2E4D4702B0C3E38B35&CFID=550559368&CFTOKEN=13078748&__acm__=1444205909_e1b0402c13a974f603db2271da839e99)